



The Ejoshi AI Starter Kit

*A calm way to start using AI for
school, life, and learning.*

Made for the Ones Who Move Different.

— EJOSHI —

v1.4 · ejoshi.ai

SECTION 01

WELCOME

You're not late.

You're not behind. You're early — to a kind of skill most schools haven't figured out how to teach yet.

This kit is the calmest, shortest way to start using AI without getting overwhelmed, without falling for hype, and without doing anything you'd be embarrassed to put your name on.

It fits in about ten minutes a day for a week. No upsell. No tricks. No "you'll get left behind." No fake urgency. Just the small, real first moves that turn AI from a thing you've heard about into a tool you can actually use.

If you've tried to learn AI before and bounced off — same. Most of what's out there is either too hyped, too technical, or too vague. This one is none of those. It assumes you're smart, you're busy, and you'd rather have one useful thing in your hand than a list of fifty bookmarks you'll never open.

Who this is for

- Anyone learning AI who wants to be smart about it, not sneaky.
- Parents who want to understand what their kids are actually doing.
- Teachers who want a calm way to bring AI into the room.
- Anyone who's been told they should "learn AI" and doesn't know where to start.

The Ejoshi promise

Always honest. Built for everyone, not just the people who already know.

Free to access. No paywall. No gatekeeping. The future shouldn't be locked behind a paywall.

What's inside. Three questions to ask before any prompt. A seven-day ramp you can start tomorrow. A short playbook for when AI gets it wrong. A plain-language glossary. Three quick adaptations for school, family, and work. Ten minutes a day. One week. That's it.

SECTION 02

BEFORE YOU TYPE

The three questions to ask before you use AI for anything.

Before you type a prompt — any prompt — ask yourself these three. They take ten seconds and save you from most of the trouble people get into with AI.

1. What am I trying to learn?

Not what answer you want. What understanding you want by the end. "Solve this problem" is an answer. "Get better at this kind of problem" is learning. AI is good at both, but only if you tell it which.

Example. Weak: *"Solve this math problem: $2x + 7 = 15$."* Strong: *"Walk me through how to solve $2x + 7 = 15$ the way you'd explain it to someone new to algebra, so I can do the next one myself."* The second prompt gets you the same answer plus the skill to find it again.

2. Would I be okay with my teacher (or boss, or kid) seeing this prompt?

If yes, you're probably fine. If you'd want to delete the chat — that's the gut check telling you something. AI is a tool. Tools work best when you use them where you can stand by what you did.

Example. Weak: *"Write my essay about the American Revolution."* Strong: *"Here's my essay draft about the American Revolution. Read it and tell me which paragraph is weakest, what's missing, and one stronger opening sentence."* One you'd hide. One you'd share.

3. Did I try first?

AI is best when it helps you finish what you started, not when it does what you haven't tried. Five minutes of your own thinking before you ask makes the AI's answer ten times more useful. You'll know what to keep and what to throw out. You'll catch where it gets things wrong.

Example. Weak: *"Plan my week for me."* Strong: *"Here's my week — three deadlines, two appointments, and a project I keep avoiding. I tried sketching a schedule but it doesn't feel right. What am I missing?"*

That's it. Three questions. They're a habit, not a rule — the more you use them, the less you'll need them.

Why these three (and not ten)

Most "how to use AI" guides give you a long checklist. That's the problem. Long checklists don't survive the moment when you're tired, in a hurry, and just want an answer.

Three questions fit in ten seconds. They handle most situations. The first one points the AI at the right job — learning vs. answering. The second one keeps you honest about whether you'd own what you're asking for. The third one keeps the AI from doing your thinking for you.

If you internalize only these three, you've got a stronger AI practice than most people who've taken a course.

SECTION 03

YOUR FIRST 7 DAYS WITH AI

A short on-ramp. One small thing each day. No streaks, no scoreboards — just a week of low-stakes practice so the tool starts to feel like yours.

Day 1 · Ask it to explain something you already half-know

Pick something from work, school, or daily life — a term, a concept, a process. Ask the AI to explain it like you're new to it. Notice what it gets right and where it skips a step.

The point isn't the explanation. The point is calibration. You're learning how the AI talks, where its examples come from, and which parts feel real versus which feel generic. That calibration is what makes you better at every prompt that follows.

Try this: *"Explain how interest on a savings account actually works, like I'm new to it. Use an everyday example."*

Day 2 · Rewrite one sentence

Take a sentence you've written — a text, an email, a caption — and ask the AI for three different versions: more confident, more friendly, shorter. Pick the one closest to how you'd want to sound.

Try this: *"Rewrite this in three versions — more confident, more friendly, and shorter: [your sentence]."*

Day 3 · Compare two options

Anything you're deciding between: two recipes, two phone plans, two approaches to a project. Give the AI the trade-offs and ask for a clear comparison.

Try this: *"I'm trying to decide between A and B. Here's my situation: [context]. What are the trade-offs, and which fits better if I care most about [priority]?"*

Day 4 · Catch it being wrong

Ask the AI a question you already know the answer to — something local, something recent, something in a language you speak. See where it gets things wrong. This is the most useful exercise in the kit. It builds the

habit of not trusting it blindly.

Try this: *"Tell me about [your hometown or recent local event]. Give me three specific facts."* Then check what it said against what you know.

Most people skip this day. Don't. The first time you catch AI confidently inventing a "fact" is the moment you stop treating it like a search engine and start treating it like a tool that needs a human in the loop. That shift is worth more than every other day in this kit combined.

Day 5 · Give it a role

The same question gets a different answer when the AI knows who to be. Try giving it a role: a patient teacher, a tough editor, a kind older sibling. Watch how the tone changes.

Try this: *"You are a patient teacher explaining this to a curious kid. Explain why the sky is blue."*

Day 6 · Use it to think, not to finish

Pick something you're stuck on — a decision, a project, a hard email — and use the AI to think it through, not to write it for you. Ask it to challenge your assumptions or point out what you might be missing.

Try this: *"I'm stuck on [problem]. Here's how I'm thinking about it: [your thinking]. What am I missing? What's a question I haven't asked yet?"*

Day 7 · Save your best prompt

Look back at the six you tried. Pick the one that gave you the most useful answer. Save it in a notes app. Customize the words that change (the topic, the role, the constraint) and keep the structure. That's your first reusable AI prompt — and the start of how every AI-fluent person works.

The point. Seven days, seven moves. You didn't try to learn AI. You used it. That's the difference.

SECTION 04

WHEN AI GETS IT WRONG

It will. That's the most important sentence in this kit.

AI doesn't know what it doesn't know. It speaks confidently about everything, including things it's just making up. That's not a flaw to fear. It's a feature to plan for.

Here's how it usually goes wrong, and what to do.

The four common wrongs

1. It makes up facts

Names, dates, quotes, citations — all sound real, none are. The clue: it's specific without being checkable.

What to do: verify anything you'd use publicly. Quick search. Trusted source. Two sources better than one.

2. It's out of date

Most AI tools have a training cutoff. They don't know what happened after that. They don't always say so.

What to do: if it matters that the answer is current, ask the AI when its knowledge ends — and check current sources for anything time-sensitive.

3. It misses local context

Something true broadly may be wrong for your town, your school, your country. AI defaults to averages.

What to do: add the local context explicitly. "I'm in [place]. Does this apply here?"

4. It agrees with you when you're wrong

AI often mirrors your framing. If you ask "why is X bad?", it'll list reasons X is bad — even if X is fine. **What to do:** ask the opposite question. "What are three reasons this might be wrong?" or "What would someone who disagrees say?"

The verification habit. Before you act on an AI answer that matters, ask one question: *how would I know if it were wrong?* If you can't answer that, check before you act.

How to know when to verify

Not every answer needs a fact-check. Asking the AI to brainstorm party themes? Run with what it gives you. Asking it for a quote, a medication dose, a legal rule, or a date you'll cite later? Verify.

A simple test: *would I be embarrassed if this turned out to be wrong?* If yes — verify. If no — use it as a starting point. The middle ground (most things) is "trust but check the load-bearing parts." Names, numbers, citations, claims about real events. Those are the load-bearing parts. The rest is texture.

The real upgrade. When you can spot AI being wrong, you've already passed the people who use it without questioning it. That's the move.

SECTION 05

THE PLAIN-LANGUAGE AI GLOSSARY

A short list of words that show up everywhere when people talk about AI. Each one explained in one sentence. No jargon.

- **Prompt** — the question or instruction you give the AI. The more specific, the better the answer.
- **Model** — the AI itself. Different models have different strengths.
- **Chat / conversation** — a back-and-forth with the AI, where it remembers what you said earlier in that session.
- **Context window** — how much the AI can hold in its memory at once during a chat. Goes back to zero in a new chat.
- **Hallucination** — when AI invents something that sounds true but isn't. The reason verification matters.
- **Training data** — everything the AI learned from. Usually a giant pile of internet text up to a cutoff date.
- **Cutoff date** — the last day the AI's training data includes. It doesn't know what happened after.
- **System prompt** — invisible instructions the AI follows in the background. Custom assistants use these to set tone and rules.
- **Fine-tuning** — extra training to make an AI better at a specific job, like coding or summarizing.
- **Token** — small pieces of text (often part of a word) that AI reads and writes in. Why long prompts can get cut off.
- **Vibe check** — not technical. Does the answer feel right? Does it match what you know? That's a real check, not a joke.
- **Verification** — confirming what AI told you, before you act on it. The most important habit on this list.
- **Open access** — free to use, no paywall, no signup, no gatekeeping. Most leading AI tools have a free tier.
- **Inference** — what's happening when you press send. The AI is running the question through the model and writing back. Most tools call this "thinking."
- **Multimodal** — an AI that can handle more than just text. Some can read images, listen to audio, watch video, or generate pictures.
- **Agent** — a step beyond a chatbot. An agent can take actions for you — browsing, scheduling, sending — usually with your permission. Use with care.

You don't need to memorize this. You'll pick these up as they come up. Keep this page nearby for the first month.

SECTION 06

SCHOOL · FAMILY · WORK

Three small adaptations. Same kit, three doorways.

At school

Use AI to learn, not to skip the learning. If you'd be okay handing the prompt to your teacher, you're probably okay. The line: AI helps you understand; you do the understanding.

A useful prompt for any subject: *"Explain [topic] like I'm new to it. Then ask me three questions to check if I got it."* The AI becomes a patient tutor you can interrupt without feeling rude.

At home (with family)

AI is already in the house — autocorrect, voice assistants, recommendation feeds. The question isn't whether your family uses it; it's whether you use it on purpose.

A useful prompt for any family moment: *"We're trying to plan [a weekend / a difficult conversation / a budget choice]. Here's our situation: [details]. Give us three options that fit our values."* The AI doesn't decide. It helps you see the shape of the decision.

At work

AI is best as a first-draft writer, a thinking partner, and a research assistant. It's worst when you let it have the final word. Your judgment is the value you bring.

A useful prompt for any work task: *"I'm writing [a thing] for [an audience] to achieve [a goal]. Here's my voice: [direct / warm / formal]. Write a first draft under [length]. I'll edit."* Then you edit until it sounds like you.

The same pattern works for research, planning, and communications. Give the AI the context, the goal, the audience, the voice, and the length. Then edit. The AI doesn't know who you are or who you're writing to. You do. The edit is where the value lives.

A note on confidential work: AI tools log most conversations. Some have stricter privacy settings; check before pasting anything you wouldn't put in an email. Client names, internal financials, and any data covered by a confidentiality agreement should stay out of public AI tools unless your employer has cleared an enterprise version.

The pattern. AI doesn't change what you do. It changes how much of it you can do, and how fast you can think through it. The skill — your skill — is still yours.

One small habit that beats most courses

People who get good at AI fast aren't smarter. They just notice. Notice when a prompt worked. Notice when an answer was useful. Notice when the AI agreed with you too quickly. Save the prompts that landed. Throw out the ones that didn't.

Ten minutes a week of noticing beats five hours of reading about AI. The tool teaches you if you let it.

What this kit does not promise

It won't make you an AI expert. It won't replace a class. It won't get you a job. It won't make you faster overnight.

What it will do — if you spend the week — is move you from "AI is a thing other people use" to "AI is a tool I know how to handle." That's the shift everything else builds on. Once you have it, the rest is just practice.

The pace. This kit is meant to be slow on purpose. Three pages on Day 1. Three pages on Day 2. You're not behind if you read it over two weeks instead of one. The point isn't speed. The point is the habit that sticks after.

SECTION 07

THE PROMISE BEHIND THIS KIT

Ejoshi is named after two sons and a family name. A father's decision to make sure his kids — and every family like his — face the AI era with a clear way to learn.

This kit is the first step of that promise, in your hands. Five pages, no upsell, no tricks. Made because there should be at least one place online where you can learn AI calmly, without being sold to, without being scared, without being talked down to.

If this helped, the next step is at ejoshi.ai/resources — more kits, all free to access, all built the same way. None of them require an account. None of them ask for a credit card. The future shouldn't be locked behind a paywall.

Welcome in. You're early. That's the best place to be.

This kit ships free for the same reason it was made calm: the people who need this most are the people who'd be most put off by a paywall, a popup, or a pitch. So there isn't one. Share the kit with anyone who'd find it useful. Print it. Translate it. Use it.

When you're ready for the next piece — there's a whole shelf of these being built. Each one made the same way. Each one free.

Start with Day 1. Save this PDF. Come back when you need it.

— Ejoshi

ejoshi.ai