



The First AI Lesson

*A classroom activity that
builds the AI verification habit.*

Made for the Ones Who Move Different.

— EJOSHI —

v1.4 · ejoshi.ai

SECTION 01

WHY THIS LESSON FIRST

Before students learn to use AI well, they need to learn to question it.

The most dangerous AI skill is confident use without critical thinking. A student who trusts AI's first answer will eventually quote one as fact. A student who's been bitten — caught AI inventing a "fact" in real time — never makes that mistake the same way again.

This is the lesson that builds the bite. One class period. Forty-five to sixty minutes. No prerequisites except internet access and one device per group.

What students leave with: not knowledge about AI. A habit. A personal verification rule, written in their own words, that they'll use the rest of their lives.

What this lesson assumes

- Your students have at least heard of AI (most have).
- You have at least one device per small group (a phone is fine).
- You can use a free AI tool in class (every leading AI has a free tier).
- You're willing to let students see AI get something wrong on purpose. That's the point.

Why this lesson works at any learning stage

The same lesson runs for students new to AI and students who use it daily. The questions change — Step 2's challenge depends on what your students already know — but the structure is identical. The dawning moment, when a student watches AI invent a fact, has the same weight at any level. That moment is the lesson. Everything else is scaffolding around it.

The Ejoshi promise. Always honest. Built for everyone, not just the people who already know. Free to access. No paywall. No gatekeeping. The future shouldn't be locked behind a paywall.

SECTION 02

PREP CHECKLIST (THE NIGHT BEFORE)

A short list. Run through it the night before, not the morning of.

- [] Device access confirmed. One device per group of 3-4. Make sure they can reach the AI tool from school WiFi.
- [] You've pre-tested three questions you know the answer to. One local, one recent, one in a language your students speak. Note where AI gets each one wrong.
- [] Whiteboard space cleared for the post-lesson list.
- [] You've decided which AI tool the class will use. Pick one. Don't have groups compare tools today — that's a different lesson.
- [] You've thought about what to do if a student says "AI was right, you're wrong." (You probably weren't — but if you were, model the correction. It's part of the lesson.)

The mindset. You are NOT teaching students that AI is bad. You're teaching them to verify before they trust. That distinction is the whole framing.

SECTION 03

STEP 1 — THE BAIT (10 MIN)

The opening. Set the stage. Let students share what they think they know.

The script

Ask: *"What do you think AI knows?"*

Take three or four answers. Don't argue with any of them. Listen.

Then ask: *"What do you think AI gets wrong?"*

This is the key moment. Most students will hesitate. They'll say "I don't know" or "Not much" or "It's pretty smart." A few will mention one thing. Almost none will have a real answer.

That hesitation is the lesson. Write on the board: *"What we don't know yet — what AI gets wrong."* And underline it.

Then say: *"Okay. Let's find out."*

What to watch for

The student who says "AI is always right." That student needs Step 2 the most. Don't single them out — but watch how they react when the AI gets it wrong in the next ten minutes.

What NOT to do

Don't tell them the answer. Don't say "actually, AI gets a lot wrong." Don't editorialize. The lesson works because they discover it. If you give them the conclusion now, the rest of the activity is theater.

SECTION 04

STEP 2 — THE TEST (15 MIN)

The exercise. Groups work; you circulate.

The setup

Groups of 3-4. Each group picks one question to ask the AI. They must already know the answer to the question. Three kinds work best:

1. Local: *"Who won the [recent local sports event]?"* or *"What was the name of the [recent local news event]?"* 2. Recent: *"What happened in [a news event from the last month]?"* 3. Native language / culture: *"What is [a word in a student's home language] in English?"* or *"What does [a local cultural reference] mean?"*

Each group asks the AI their question. They write down the AI's answer. Word for word.

The script while you circulate

Walk around. Ask each group quietly: *"What did it say? What's actually true?"*

Don't react when they catch a mistake. Just note it. Say: *"Save that. We'll come back to it."* Move on.

What you're watching for

The dawning. The moment the group realizes the AI just said something confidently wrong. That moment — usually around the 8 to 10 minute mark — is the lesson. Some groups hit it early. Some need a nudge ("are you SURE that's what happened?"). One or two groups might not catch a mistake at all on their first try. That's fine. Have them ask a second question.

A nudge that almost always works

If a group is stuck — the AI's answer sounds plausible and they aren't sure — try: *"Ask it the same question but more specific. Include a date, a name, a place."* AI tends to invent the specifics. The more you pin it down, the more often it slips. That slip is what the group is hunting for.

SECTION 05

STEP 3 — THE CATCH (10 MIN)

Groups report back. The class sees the pattern together.

The script

Each group: 30 seconds. Two questions:

1. What did AI get right? 2. What did AI get wrong — and what did it say confidently?

Write the wrongs on the board. Don't editorialize.

When all groups have shared, ask the room: *"What's the pattern? What kinds of things did AI get wrong?"*

You'll usually hear: local things, recent things, language nuance, specific names. Write the categories on the board.

Then ask the question that lands the lesson: *"How would you know if you didn't already know the answer?"*

That's the question. Let it sit. Don't fill the silence. Watch the room work.

The sentence to land

Somewhere in this step, say it out loud: *"AI is confident even when it's wrong. The only thing that protects you is knowing when to check."*

That sentence is the one students will remember six months from now.

SECTION 06

STEP 4 — BUILD THE HABIT (10 MIN)

Each student writes their personal verification rule. Not a class rule. Their own.

The prompt

On a small piece of paper, each student writes one sentence in this shape:

"When I use AI, I will always verify by _____ before I trust an answer about _____."

Examples (don't show these until they've tried):

- *"When I use AI, I will always verify by checking the date before I trust an answer about anything that happened recently."*
- *"When I use AI, I will always verify by asking my mom before I trust an answer about something in [our home language]."*
- *"When I use AI, I will always verify by checking two sources before I trust an answer that I'm going to put in an essay."*

Why the student writes their own

A rule the student wrote is a rule the student remembers. A rule the teacher wrote is a rule the student forgets. The whole exercise — Steps 1 through 4 — was building to this sentence. The sentence is the takeaway. Everything else was preparation.

One pattern to notice

The rules students write usually reveal what they care about. The student who writes "I'll check before I trust an answer about my family's culture" is telling you something. The student who writes "I'll check before I trust an answer in my college essay" is telling you something different. Both rules are good. Both reveal the student. Read them with that in mind.

What to do with the sentences

Pin them up. Or collect and return them next class. Or have students keep them in a notebook they'll see again. Anything but throw them away. The rule needs a place to live.

SECTION 07

WHAT THIS LESSON ISN'T

A few things students sometimes assume after the lesson — that aren't the takeaway.

This lesson isn't "AI is bad." It's "AI is a tool that requires you to think." Those are different. A student who walks out thinking AI is bad will avoid it; a student who walks out thinking AI requires thought will use it well.

This lesson isn't "always check everything." That's exhausting and unsustainable. The takeaway is more useful: check the things that matter. Names, numbers, citations, claims about real events, anything you'd be embarrassed to be wrong about. The rest can stay light.

This lesson isn't a one-time event. It's a foundation. The verification habit gets reinforced every time the student notices AI being wrong in any future assignment. The single lesson plants the seed. Your reinforcement over the year is what grows it.

The takeaway behind the takeaway. "Verify the important parts" travels far beyond AI. The student who learns it here will use the same habit with news headlines, social media claims, and confident-sounding strangers for the rest of their life.

SECTION 08

DEBRIEF AND ASSESSMENT

Three minutes of discussion. Not a quiz. A check on what landed.

The debrief questions

- *"What's the difference between AI being wrong and AI lying?"*
- *"Does a tool being wrong make it useless? Why or why not?"*
- *"Who is responsible for what you produce, when AI helped?"*

These don't have right answers. They're for surfacing what each student is making of the lesson.

How to assess what students learned

The traditional assessment: a quiz. Don't.

A better assessment: a week later, ask casually, *"Did anyone use AI this week? Did you verify anything?"* Listen to the answers. The students who internalized the lesson use the language of verification naturally. The ones who didn't, didn't.

Even better: the next time you give an assignment that allows AI, require students to write one sentence at the top about how they verified the AI's contributions. That sentence is the assessment. It's also the lesson, repeating.

One last assessment move

Save the verification rules students wrote in Step 4. At the end of the term, hand them back. Ask each student: did your rule hold up? Did you change it? What did you learn? That five-minute reflection at the end of the term turns the one-class lesson into a year-long arc — and reinforces the habit at the moment students are most likely to forget it.

SECTION 09

WHEN STUDENTS PUSH BACK

They will. Here are the three common kinds, and what each one means.

"AI is always right."

Means: they haven't been caught yet. The activity itself usually handles this — they'll catch it during Step 2. If they don't, give them a question you know the AI fails on, and let them ask it.

"This is dumb. Why do we need to know this?"

Means: they're testing whether you'll get defensive. Don't. Say: *"Fair question. Five years from now, when you use AI for something that actually matters — a college essay, a research paper, a work email — this is the habit that protects you. We're learning it cheap, now, instead of expensively, later."*

"My teacher last year said AI is fine."

Means: there's a real disagreement among the adults in their life. Don't dismiss the other teacher. Say: *"There are different opinions on this. I'm teaching you a habit. The habit works whether you use AI a lot or a little, whether your other teachers love it or hate it. The habit is yours."*

The pattern. Pushback isn't resistance. It's the student trying to figure out what frame to use. Give them yours, calmly, and most of them adopt it.

SECTION 10

FOLLOW-UP LESSONS

What to teach next. Three natural sequences after this lesson lands.

Follow-up 1 — The "compare answers" lesson

Same question, different prompts. Different students. Compare how the answer changes. The takeaway: the prompt shapes the answer. AI is a mirror — what you ask it shapes what you see.

Follow-up 2 — The "build the prompt" lesson

Students design prompts together. The class watches what changes when context, role, and format are added. The takeaway: the skill is in the asking.

Follow-up 3 — The "AI in this subject" lesson

What's AI good at, in YOUR subject? What's it bad at? Where does it help; where does it shortcut the learning? This conversation goes deeper after students have practiced verification and prompting.

Follow-up 4 — The "AI on this specific assignment" check-in

Two weeks before any major assignment that allows AI use, run a 10-minute check-in: where could AI help on this? Where might it hurt? What would you ask it that you'd want to verify before turning in? This rehearsal makes the assignment itself shorter and the work cleaner — and it surfaces problems before they're costly.

The arc. Verification first. Prompting second. Subject-specific use third. In that order, you build students who use AI as a tool. Skip the first step and you build students who use AI as an oracle.

SECTION 11

A NOTE FOR THE TEACHER

If you've never run a lesson like this, you might worry about losing control of the room. You won't. Students are surprisingly serious about catching the AI being wrong. The activity has a built-in dignity — they're doing real work, finding real mistakes.

If the lesson doesn't go the way you planned, that's also information. Adjust. Try it again next month with a different group. The lesson is yours now.

A practical timing note

You don't have to teach this lesson on the first day of the school year. Some teachers run it during the second or third week, once routines are settled. The "first AI lesson" name refers to the first lesson of AI — not necessarily the first lesson of the year.

When this lesson lands well

You'll know. The next time you mention AI in class, students will say "but how do we know it's right?" without being prompted. That sentence is the lesson. Six months from now, it's the habit. Years from now, it's just how they think.

You're not teaching them about AI. You're teaching them to think clearly about any source that sounds confident. The lesson outlives the tool.

Start with one class. Save what worked. Come back when you need it.

— Ejoshi ejoshi.ai