How to Read About AI

Navigate the noise
without losing your footing.

Made for the Ones Who Move Different.

— EJOSHI —

v1.4 · ejoshi.ai

SECTION 01

THE PROBLEM WITH AI NEWS

AI generates more headlines than almost any other topic right now. Most of them are written to provoke a reaction — fear, excitement, urgency. Almost none of them are written to help you understand what's actually happening.

This kit isn't about which AI is best. It's about how to read about AI without getting played.

Inside: the three kinds of AI stories that dominate the feed and what each one is doing. Five questions to ask before you react. A short decoder for company press releases. A hype lexicon — the words that should slow you down. A 30-second test for "should I forward this?" And a short list of sources that are usually careful.

The whole kit fits in about ten minutes. The skills last a lifetime. Most of them aren't AI-specific. They're media-literacy skills that happen to apply to AI right now.

Who this is for

- Anyone trying to follow AI news without getting hyped or scared.
- Parents whose kids are sending them "AI is going to take over the world" links.
- Professionals who need to know what's real in AI for their work.
- Anyone who's clicked an AI headline and felt worse afterward.

A short note on why this kit exists

Most "AI literacy" guides teach you about AI. This one teaches you to read about AI. Those are different skills. You don't need to understand transformer architecture to spot a bad article. You need a small set of habits that filter the signal from the noise. That's the whole kit.

> **The Ejoshi promise.** Always honest. Built for everyone, not just the people who already know. Free to access. No paywall. No gatekeeping. The future shouldn't be locked behind a paywall.

SECTION 02

THE THREE TYPES OF AI STORIES

Most AI news falls into one of three buckets. Knowing which bucket a story is in tells you how to read it.

Type 1 — Benchmark stories

The shape: *"AI beats humans at [task]."*

What they are: a research lab or company tested their AI on a specific task and beat a previous score. Usually true, often misleading.

The trap: the task was narrow. The result doesn't mean what the headline implies. "AI passed the bar exam" doesn't mean AI is now a lawyer. It means AI did well on a specific multiple-choice test that measures something — but not the actual practice of law.

What to ask: *Does this capability change anything I do this week?* If yes, look into it. If no, file under "interesting but not urgent."

Worked example: a headline about AI passing a medical licensing exam. Useful to know AI is good at structured medical knowledge. Useful for thinking about future tools. Useless as a reason to trust AI with your actual health questions today.

Type 2 — Prediction stories

The shape: *"AI will [do something dramatic] by [year]."*

What they are: speculation by someone who may or may not have any track record on prediction. Often a CEO, sometimes a researcher, occasionally a pundit.

The trap: predictions sound like facts. They're not. They're guesses, often by people with a financial interest in the prediction being taken seriously.

What to ask: *Who is making this prediction, what do they benefit from people believing it, and have they been right about anything before?*

Worked example: "AI will replace 50% of office jobs by 2030." Mark the date. Save the article. Check back in 2030. Most of these predictions don't survive contact with the actual decade they predicted.

Type 3 — Application stories

The shape: *"[Company] uses AI to [do specific thing]."*

What they are: a real use case from a real organization. Sometimes vendor-promoted, sometimes journalist-discovered.

The trap: smaller than the others. The main risk is overgeneralizing — assuming if it works for them, it'll work for you.

What to ask: *What did they actually do, what was the result, and what was the cost?*

Worked example: a small law firm using AI to draft first-pass contracts. Useful to read because it's concrete — they did X, it saved Y hours, they kept Z under human review. You can adapt this to your own situation in a way you can't with the other two types.

The bonus type — failure stories

Worth noting: AI failure stories are the most underrated category. Companies and writers don't promote them as much, but they're the most teachable. *"We tried AI to do X. Here's what went wrong. Here's what we'd do differently."* When you see one, read it carefully. The mistakes other people document are the ones you don't have to make yourself.

> **The shortcut.** Benchmark stories — interesting, file away. Prediction stories — track but don't act on. Application stories — the most useful, especially for your specific work.

SECTION 03

FIVE QUESTIONS TO ASK BEFORE YOU REACT

Run these on any AI story that's making you feel something — hyped, scared, urgent, or amazed.

1. Who wrote this, and what do they get out of your reaction?

Newsrooms want clicks. Companies want investment. Researchers want grants. Influencers want followers. None of those incentives are evil — but they shape what gets written. The question to ask: if I were the writer, and I needed [outcome], what kind of story would I write?

2. Is this a study result, a product announcement, or an opinion?

These read similarly but mean different things. A study result is usually peer-reviewed and replicable; a product announcement is marketing; an opinion is one person's read. Headlines blur all three. The body of the article usually tells you which it is.

3. What does the headline claim versus what does the article actually say?

Headlines and articles often disagree. A headline saying "AI threatens X" might be supported by an article that ends with "though this would require Y, which is unlikely." Read past the headline. Always.

4. How big was the sample, and how controlled was the test?

For benchmark and study stories: if the test was small or the conditions were narrow, the result might not transfer. AI passing a test on one dataset doesn't mean AI passes that test in the wild. Sample size and conditions are buried in the article — usually. If they're not there at all, that's information too.

5. What would the opposite headline look like — and is it also true?

For most AI claims, the opposite is also true at the same time. *"AI is the most disruptive technology in history"* and *"AI is overhyped and most use cases are mundane"* are both defensible right now. The interesting work is figuring out where the truth lives between them.

> **Why these five.** They're media-literacy questions, not AI questions. They work on any topic that generates excitement. AI just happens to be the noisiest one right now.

SECTION 04

THE COMPANY PRESS-RELEASE DECODER

A press release is the company telling you what they want you to think. Reading one well takes practice.

What to focus on

The three sentences that matter most in any AI product announcement:

1. **What specifically can it do today** — not "in the future," not "soon," not "we believe." Today, in version 1. 2. **What it costs** — including hidden costs (data, time, integration, lock-in). If the price isn't named, the announcement is incomplete. 3. **What it doesn't do** — usually buried. Sometimes missing entirely. The absence of this section is itself information.

What to skim past

- The opening paragraph about how this changes everything.
- Quotes from executives saying it's revolutionary.
- Generic statistics about the size of the market.
- Anything in the future tense without a specific date.

The decoder test

After reading the release, ask: *"If I had to summarize what's new in one specific sentence, could I?"* If yes, the release was substantive. If no, the release was mostly wallpaper.

The "wait six months" rule

The single most useful habit with company press releases: don't react in the first week. Wait six months. By then, you'll know if the thing actually shipped, if it actually works, and if the use cases the company predicted turned out to be the real use cases. Most "revolutionary" AI products either quietly underdeliver or actually deliver in a different shape than promised. Six months gives the dust time to settle.

> **The honest catch.** Most press releases are mostly wallpaper. That's not a moral failing — it's the job of a press release. Your job is to find the one sentence of substance inside it.

SECTION 05

THE AI HYPE LEXICON

Words that should slow you down when you see them in an AI article. None of these are automatically lies. They're just signals — when you see one, read more carefully.

- **Revolutionary** — Sometimes true. Usually marketing. Ask: revolutionary compared to what specifically?
- **Unprecedented** — Almost never literally true. Usually means "the company wants you to feel like this is bigger than it is."
- **Game-changing** — A claim that the rules of some game just changed. Ask: which game, and how exactly are the rules different?
- **AGI / superintelligence** — Loaded terms that mean different things to different people. When you see one, find out which definition the writer is using before you continue.
- **Sentient** — Almost certainly false in any current AI. AI is good at sounding like it's thinking; that's a feature of large language models, not evidence of consciousness.
- **Will replace [job]** — A prediction in disguise. Translate as: *"someone wants you to believe this."*
- **Self-aware** — Even more loaded than sentient. Usually used by writers who haven't thought carefully about what self-awareness means.
- **Solves [problem]** — A strong claim. Ask: does it solve the problem, or does it improve a partial solution? Almost always the latter.

A trick for noticing hype

Read the article out loud. Specifically, read the sentences that use these loaded words out loud. Hype words almost always come with a tone you can hear when you're forced to say them. The same sentence read silently sounds reasonable; read aloud, it sounds like marketing. Try this on the next AI article that's making you feel something. Your own voice is a hype detector.

> **What this list isn't.** A blocklist. You don't ignore articles that use these words. You just read those articles more skeptically — and look for the specific claims underneath the loaded words.

SECTION 06

BEFORE YOU FORWARD — A 30-SECOND GUT CHECK

Two questions to ask before you share any AI article in a group chat, on social, or with a worried family member.

Question 1 — Have I read past the headline?

If the answer is no, don't forward yet. Read first. Forward second.

This is the single biggest filter. Most of the worst AI takes circulating online are from people who shared a headline without reading the article. The article often softens the claim or names a constraint that changes everything. The headline is the bait; the article is the truth, often.

Question 2 — Would I bet money this claim is true a year from now?

Most AI headlines wouldn't survive that bet. The forecast in the article will turn out to be partly right, partly wrong, partly irrelevant. If you wouldn't bet on it, you probably shouldn't share it as if you would.

A bonus question — am I sharing this to inform, to vent, or to bond?

All three are fine reasons to share things. But they shape how the message lands. If you're venting, say so. If you're bonding over a shared worry, name it. If you're informing, make sure you've actually checked what you're sharing.

> **The honest catch.** This isn't a rule against forwarding things. It's a 30-second pause that catches the worst forwards before they leave your phone. Most of what doesn't survive the gut check shouldn't have been shared anyway.

SECTION 07

SOURCES WORTH YOUR TIME

No source is perfect. These are usually more careful than the median.

- **MIT Technology Review** — Longer, more nuanced pieces. Useful for understanding why a technology works, not just that it does.
- **The Gradient** — Research-adjacent, technical but readable. Useful when you want depth from people doing the work.
- **Import AI (newsletter)** — A weekly written by someone close to the research. Useful for tracking what experts are paying attention to.
- **Ars Technica** — Strong on how things actually work mechanically. Useful when you want to understand the engineering, not just the implications.

What you can safely ignore

- Headlines containing BOMBSHELL, TERRIFYING, GAME CHANGER, or a countdown.
- AI sentience claims from non-peer-reviewed sources.
- Specific-year predictions from people who've been wrong before.
- Company press releases reprinted as news with no independent reporting.
- "AI is going to take all jobs by [near year]" — they've been saying this since the 1960s. The economy is messier than that.

SECTION 08

THE ONE-LINE TEST

If you can only remember one thing from this kit, remember this:

> *"If a headline made me feel something strong, that's the moment to slow down — not speed up."*

Strong feelings about a headline are the goal of the headline — that's literally how they're written, by people whose job it is to provoke them in you. Strong feelings turned into a forward, a comment, or a decision are usually a mistake. The pause is the skill.

Most AI literacy is built in the thirty seconds between reading a headline and reacting to it. Use that thirty seconds. The headlines will keep coming. Your attention is the scarce resource.

The bigger picture

AI is the loudest beat in the news right now, but it isn't the first. Every new technology has been covered like this — telephones, electricity, cars, the internet, smartphones, social media. Each one was going to change everything, end work as we know it, and either save the world or end it. Some of those predictions were closer than others. None of them were as dramatic as the headlines.

That's the long view AI deserves. It will change things. It already has. The changes will be smaller than the loudest predictions and bigger than the quietest ones. The truth, as usual, will live between them. Your job — and the job of this kit — is to be the kind of reader who can find that middle, calmly, while everyone else is shouting.

Start with one headline. Save the questions that helped you read it better. Come back when you need them.

— Ejoshi ejoshi.ai